

Parallel Computer Architecture: Homework 1

February 1, 2012

1 2.7

Problem description:

- 2.7 Gaussian elimination is a well-known technique for solving simultaneous linear systems of equations. Variables are eliminated one by one until there is only one left, and then the discovered values of variables are back-substituted to obtain the values of other variables. In practice, the coefficients of the unknowns in the equation system are represented as a matrix A , and the matrix is first converted to an upper-triangular matrix (a matrix in which all elements below the main diagonal are 0). Then back-substitution is used. Let us focus on the conversion to an upper-triangular matrix by successive variable elimination. Pseudocode for sequential Gaussian elimination is shown in Figure 2.18. The diagonal element for a particular iteration of the k loop is called the *pivot element*, and its row is called the *pivot row*.
- Draw a simple figure illustrating the dependences among matrix elements.
 - Assuming a decomposition into rows and an assignment into blocks of contiguous rows, write a shared address space **parallel** version using the primitives used for the equation solver in this chapter.
 - Write a message-passing version for the same decomposition and assignment, first using synchronous message passing and then any form of asynchronous message passing.
 - Can you see obvious performance problems with this partitioning? (We will discuss this further in the next chapter.)
 - Modify both the shared address space and message-passing versions to use an interleaved assignment of rows to processes.
 - Discuss the trade-offs (programming difficulty and any likely major performance differences) in programming the shared address space and message-passing versions.

```

procedure Eliminate (A)      /*triangularize the matrix A*/
begin
  for k ← 0 to n-1 do      /*loop over all diagonal (pivot) elements*/
    begin
      for j ← k+1 to n-1 do /*for all elements in the row of, and to the right of,
                             the pivot element*/
         $A_{k,j} = A_{k,j} / A_{k,k};$  /*divide by pivot element*/
       $A_{k,k} = 1;$ 
      for i ← k+1 to n-1 do /*for all rows below the pivot row*/
        for j ← k+1 to n-1 do /*for all elements in the row*/
           $A_{i,j} = A_{i,j} - A_{i,k} * A_{k,j};$ 
        endfor
         $A_{i,k} = 0;$ 
      endfor
    endfor
  endfor
end procedure

```

FIGURE 2.18 Pseudocode describing sequential Gaussian elimination

Figure 1: figure 2.18

1.1 2.7(a)

There are three cases as showed following. Red is used to mark the target entry. Shadow area shows the dependent entries in the matrix.




Figure 2: dependences of entry $a_{i,j}$, where $i < j$




Figure 3: dependences of entry $a_{i,j}$, where $i \leq j$

In the last case, we have $a_{i,j}$ for $i = j$.




Figure 4: dependences of entry $a_{i,j}$, where $i = j$

1.2 2.7(b)

See `src/b.cpp`.

1.3 2.7(c)

TODO: what's any form of async message passing?

Here's sync message passing program. (draft code, not compiled)
See src/c.cpp

1.4 2.7(d)

The only issue I see about the partitioning is that each pivot has different number of workload. This imbalance could cause some threads/processes remaining idle.

1.5 2.7(e)

I didn't see a reason why we want an interleaved version. There's no dependency between consecutive rows.

1.6 2.7(f)

In this particular problem, I would prefer to use shared variable approach.

1. It's much shorter, as I mentioned in the class.
2. Message passing version involves a huge data communication cost,

especially in a network environment. Because in each iteration, we need to spread out almost the whole matrix.